


# Preference Learning: A Tutorial Introduction

**Johannes Fürnkranz**

Knowledge Engineering  
Dept. of Computer Science  
Technical University Darmstadt, Germany



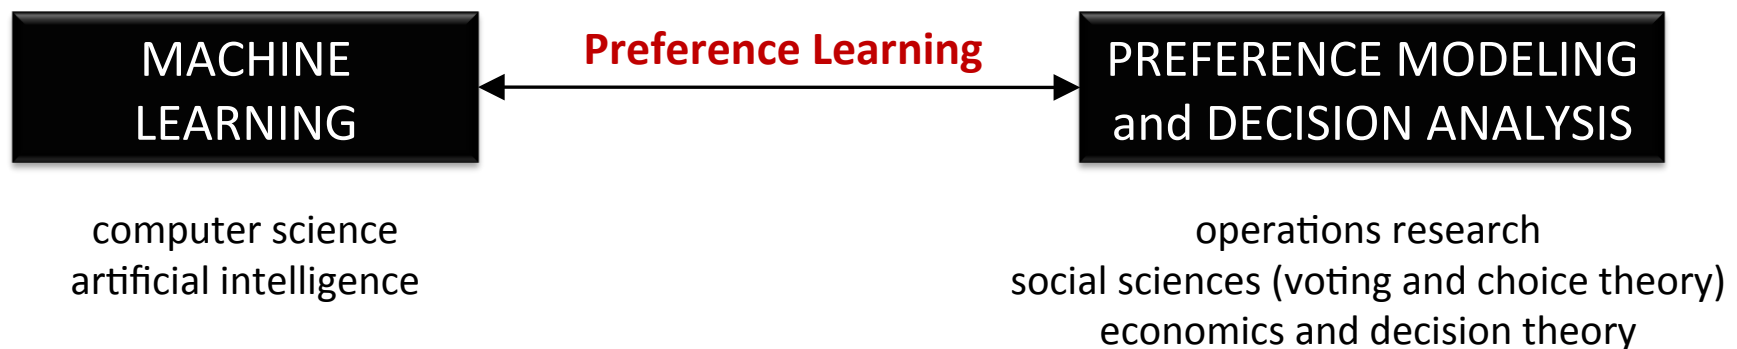
**Eyke Hüllermeier**

Knowledge Engineering & Bioinformatics Lab  
Dept. of Mathematics and Computer Science  
Marburg University, Germany



# What is Preference Learning ?

- Preference learning is an emerging **subfield of machine learning**
- Roughly speaking, it deals with the **learning of (predictive) preference models** from observed (or extracted) preference information



## Workshops and Related Events

---

- NIPS–01: New Methods for Preference Elicitation
- NIPS–02: Beyond Classification and Regression: Learning Rankings, Preferences, Equality Predicates, and Other Structures
- KI–03: Preference Learning: Models, Methods, Applications
- NIPS–04: Learning With Structured Outputs
- NIPS–05: Workshop on Learning to Rank
- IJCAI–05: Advances in Preference Handling
- SIGIR 07–10: Workshop on Learning to Rank for Information Retrieval
- **ECML/PDDK 08–10: Workshop on Preference Learning**
- NIPS–09: Workshop on Advances in Ranking
- American Institute of Mathematics Workshop in Summer 2010: The Mathematics of Ranking
- NIPS-11: Workshop on Choice Models and Preference Learning

# Preferences in Artificial Intelligence

---

More generally, „preferences“ is a key topic in current AI research

**User preferences** play a key role in various fields of application:

- recommender systems,
- adaptive user interfaces,
- adaptive retrieval systems,
- autonomous agents (electronic commerce),
- games, ...

Preferences in **AI research**:

- **preference representation** (CP nets, GAU networks, logical representations, fuzzy constraints, ...)
- **reasoning** with preferences (decision theory, constraint satisfaction, non-monotonic reasoning, ...)
- **preference acquisition** (preference elicitation, **preference learning**, ...)

# AGENDA

---

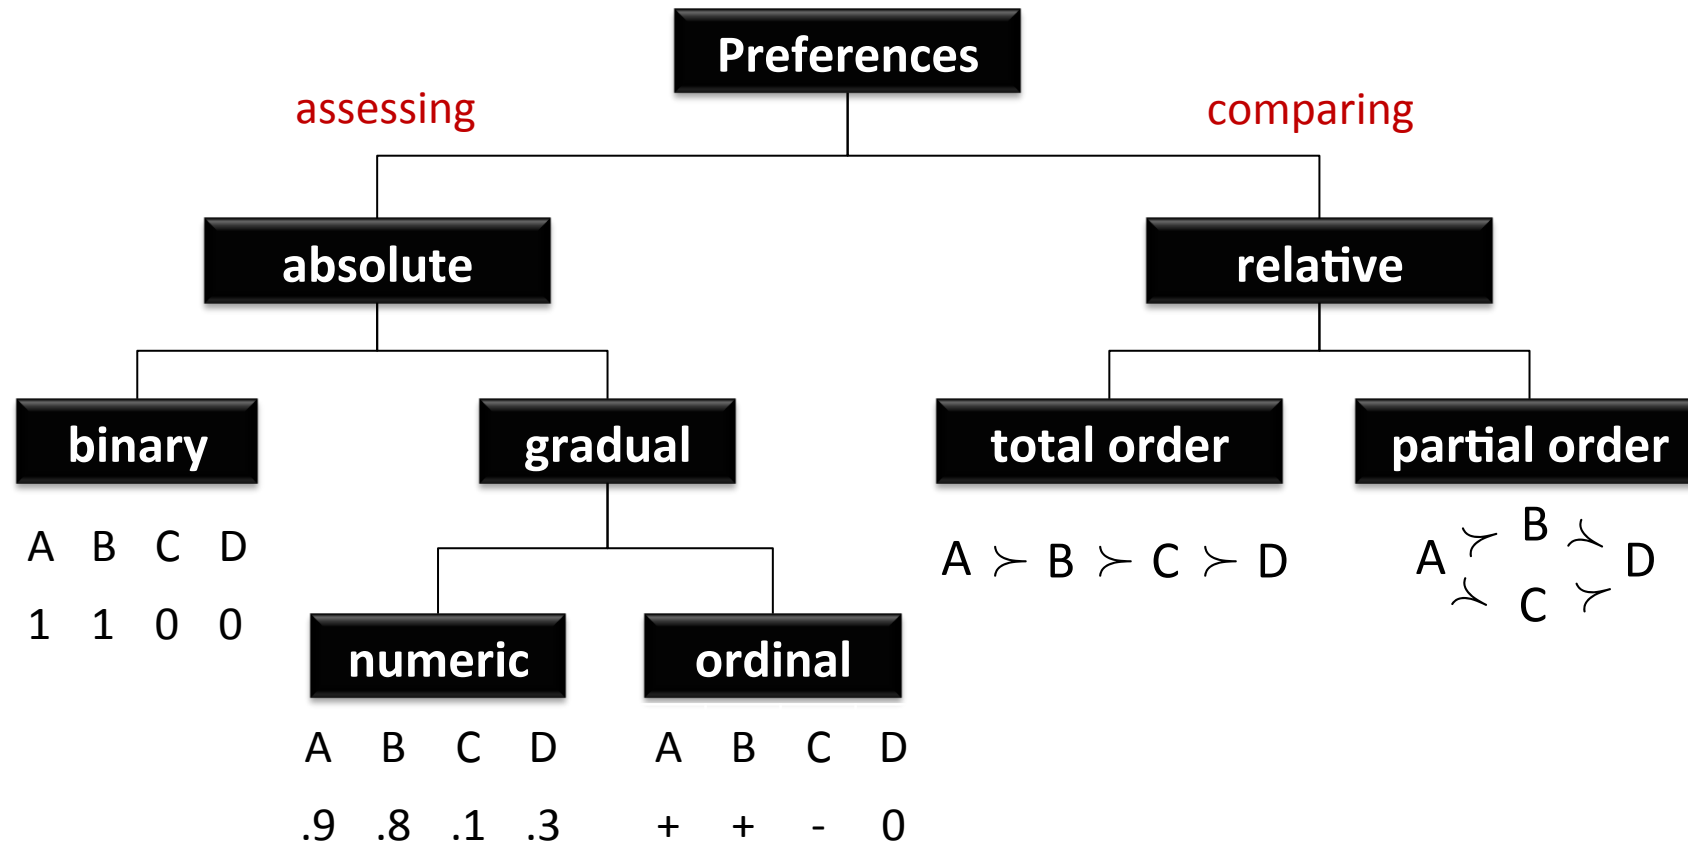
1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
3. Preference Learning Techniques
4. Conclusions

# Preference Learning

Preference learning problems can be distinguished along several **problem dimensions**, including

- **representation of preferences, type of preference model:**
  - utility function (ordinal, numeric),
  - preference relation (partial order, ranking, ...),
  - logical representation, ...
- **description of individuals/users and alternatives/items:**
  - identifier, feature vector, structured object, ...
- **type of training input:**
  - direct or indirect feedback,
  - complete or incomplete relations,
  - utilities, ...
- ...

# Preference Learning



→ (ordinal) regression

→ classification/ranking

## Structure of this Overview

---

- (1) Preference learning as an extension of **conventional supervised learning**:  
Learn a mapping

$$\mathcal{X} \rightarrow \mathfrak{P}$$

that maps instances to preference models ( $\rightarrow$  structured/complex output prediction).

- (2) Other settings (object ranking, instance ranking, CF, ...)



## Structure of this Overview

- (1) Preference learning as an extension of **conventional supervised learning**:  
Learn a mapping

$$\mathcal{X} \rightarrow \mathfrak{P}$$

that maps instances to preference models ( $\rightarrow$  structured/complex output prediction).

Instances are typically (though not necessarily) characterized in terms of a feature vector.

The output space consists of preference models over a fixed set of alternatives (classes, labels, ...) represented in terms of an identifier  
 $\rightarrow$  *extensions of multi-class classification*

# Multilabel Classification [Tsoumakas & Katakis 2007]

## Training

X1	X2	X3	X4	A	B	C	D
0.34	0	10	174	0	1	1	0
1.45	0	32	277	0	1	0	1
1.22	1	46	421	0	0	0	1
0.74	1	25	165	0	1	1	1
0.95	1	72	273	1	0	1	0
1.04	0	33	158	1	1	1	0

Binary preferences on a fixed set of items: liked or disliked

## Prediction

0.92	1	81	382	0	1	0	1
------	---	----	-----	---	---	---	---

## Ground truth

0.92	1	81	382	1	1	0	1
------	---	----	-----	---	---	---	---

LOSS



# Multilabel Ranking

## Training

X1	X2	X3	X4	A	B	C	D
0.34	0	10	174	0	1	1	0
1.45	0	32	277	0	1	0	1
1.22	1	46	421	0	0	0	1
0.74	1	25	165	0	1	1	1
0.95	1	72	273	1	0	1	0
1.04	0	33	158	1	1	1	0

Binary preferences on a fixed set of items: liked or disliked

## Prediction

				B	$\succ$	D	$\succ$	C	$\succ$	A
0.92	1	81	382	4		1		3		2

A ranking of all items

## Ground truth

0.92	1	81	382	1		1		0		1
------	---	----	-----	---	--	---	--	---	--	---



# Graded Multilabel Classification [Cheng et al. 2010]

## Training

X1	X2	X3	X4	A	B	C	D
0.34	0	10	174	--	+	++	0
1.45	0	32	277	0	++	--	+
1.22	1	46	421	--	--	0	+
0.74	1	25	165	0	+	+	++
0.95	1	72	273	+	0	++	--
1.04	0	33	158	+	+	++	--

Ordinal preferences on a fixed set of items: liked, disliked, or something in-between

## Prediction

0.92	1	81	382	--	+	0	++
------	---	----	-----	----	---	---	----

## Ground truth

0.92	1	81	382	0	++	--	+
------	---	----	-----	---	----	----	---



# Graded Multilabel Ranking

## Training

X1	X2	X3	X4	A	B	C	D
0.34	0	10	174	--	+	++	0
1.45	0	32	277	0	++	--	+
1.22	1	46	421	--	--	0	+
0.74	1	25	165	0	+	+	++
0.95	1	72	273	+	0	++	--
1.04	0	33	158	+	+	++	--

Ordinal preferences on a fixed set of items: liked, disliked, or something in-between

## Prediction

				B	$\succ$	D	$\succ$	C	$\succ$	A
0.92	1	81	382	4		1		3		2

A ranking of all items

## Ground truth

0.92	1	81	382	0	++	--	+
------	---	----	-----	---	----	----	---



# Label Ranking [Hüllermeier et al. 2008]

## Training

X1	X2	X3	X4	Preferences
0.34	0	10	174	$A \succ B, B \succ C, C \succ D$
1.45	0	32	277	$B \succ C$
1.22	1	46	421	$B \succ D, A \succ D, C \succ D, A \succ C$
0.74	1	25	165	$C \succ A, C \succ D, A \succ B$
0.95	1	72	273	$B \succ D, A \succ D$
1.04	0	33	158	$D \succ A, A \succ B, C \succ B, A \succ C$

Instances are associated with pairwise preferences between labels.

## Prediction

				B	$\succ$	D	$\succ$	C	$\succ$	A
0.92	1	81	382	4		1		3		2


A ranking of all labels

## Ground truth


0.92	1	81	382	2		1		3		4
------	---	----	-----	---	--	---	--	---	--	---



# Calibrated Label Ranking [Fürnkranz et al. 2008]



Combining absolute and relative evaluation:



## Structure of this Overview

---

- (1) Preference Learning as an extension of conventional supervised learning:

Learn a mapping

$$\mathcal{X} \rightarrow \mathfrak{P}$$

that maps instances to preference models ( $\rightarrow$  structured output prediction).

- (2) **Other settings** (no clear distinction between input/output space)

object ranking, instance ranking, collaborative filtering



# Object Ranking [Cohen et al. 99]

## Training

$(0.74, 1, 25, 165)$	$\succ$	$(0.45, 0, 35, 155)$	Pairwise preferences between objects (instances).
$(0.47, 1, 46, 183)$	$\succ$	$(0.57, 1, 61, 177)$	
$(0.25, 0, 26, 199)$	$\succ$	$(0.73, 0, 46, 185)$	
$(0.95, 0, 73, 133)$	$\succ$	$(0.25, 1, 35, 153)$	
$(0.68, 1, 55, 147)$	$\succ$	$(0.67, 0, 63, 182)$	

## Prediction (ranking a new set of objects)

$$\mathcal{Q} = \{\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \mathbf{x}_4, \mathbf{x}_5, \mathbf{x}_6, \mathbf{x}_7, \mathbf{x}_8, \mathbf{x}_9, \mathbf{x}_{10}, \mathbf{x}_{11}, \mathbf{x}_{12}, \mathbf{x}_{13}\}$$

$$\mathbf{x}_{10} \succ \mathbf{x}_4 \succ \mathbf{x}_7 \succ \mathbf{x}_1 \succ \mathbf{x}_{11} \succ \mathbf{x}_2 \succ \mathbf{x}_8 \succ \mathbf{x}_{13} \succ \mathbf{x}_9 \succ \mathbf{x}_3 \succ \mathbf{x}_{12} \succ \mathbf{x}_5 \succ \mathbf{x}_6$$

## Ground truth (ranking or top-ranking or subset of relevant objects)

$$\mathbf{x}_{11} \succ \mathbf{x}_7 \succ \mathbf{x}_4 \succ \mathbf{x}_2 \succ \mathbf{x}_{10} \succ \mathbf{x}_1 \succ \mathbf{x}_8 \succ \mathbf{x}_{13} \succ \mathbf{x}_9 \succ \mathbf{x}_{12} \succ \mathbf{x}_3 \succ \mathbf{x}_5 \succ \mathbf{x}_6$$

$$\mathbf{x}_{11} \succ \mathbf{x}_7 \succ \mathbf{x}_4 \succ \mathbf{x}_2 \succ \mathbf{x}_{10}$$

$$\mathcal{P} = \{\mathbf{x}_{11}, \mathbf{x}_7, \mathbf{x}_4, \mathbf{x}_2, \mathbf{x}_{10}, \mathbf{x}_1\} \quad \mathcal{N} = \{\mathbf{x}_8, \mathbf{x}_{13}, \mathbf{x}_9, \mathbf{x}_{12}, \mathbf{x}_3, \mathbf{x}_5, \mathbf{x}_6\}$$

# Instance Ranking [Fürnkranz et al. 2009]

## Training

	X1	X2	X3	X4	class
$x_1$	0.34	0	10	174	--
$x_2$	1.45	0	32	277	0
$x_3$	0.74	1	25	165	++
...	...	...	...	...	...
$x_n$	0.95	1	72	273	+

Absolute preferences on an ordinal scale.

## Prediction (ranking a new set of objects)

$$Q = \{x_1, x_2, x_3, x_4, x_5, x_6, x_7, x_8, x_9, x_{10}, x_{11}, x_{12}, x_{13}\}$$


$$x_{10} \succ x_4 \succ x_7 \succ x_1 \succ x_{11} \succ x_2 \succ x_8 \succ x_{13} \succ x_9 \succ x_3 \succ x_{12} \succ x_5 \succ x_6$$

## Ground truth (ordinal classes)

$x_{10}$	$x_4$	$x_7$	$x_1$	$x_{11}$	$x_2$	$x_8$	$x_{13}$	$x_9$	$x_3$	$x_{12}$	$x_5$	$x_6$
+	0	++	++	--	+	0	+	--	0	0	--	--

# Instance Ranking [Fürnkranz et al. 2009]

Extension of AUC maximization to the polytomous case, in which instances are rated on an ordinal scale such as {**bad**, **medium**, **good**}



# Collaborative Filtering [Goldberg et al. 1992]

		PRODUCTS								
		P1	P2	P3	...	P38	...	P88	P89	P90
USERS	U1	1		4	...		...		3	
	U2		2	2	...		...	1		
	...				...		...			
	U46	?	2	?	...	?	...	?	?	4
	...				...		...			
	U98	5			...		...	4		
	U99			1	...		...		2	

1: very bad, 2: bad, 3: fair, 4: good, 5: excellent

Inputs and outputs as identifiers, absolute preferences in terms of ordinal degrees.

# Dyadic Prediction [Menon & Elkan 2010]

					?	?	?	?	?	?	?	?	?	?
					10	14	45	32	52	61	16	33	53	
					P1	P2	P3	...	P38	...	P88	P89	P90	
?	?	1	5	U1	1		4	...		...		3		
?	?	0	4	U2		2	2	...		...	1			
?	?	0	6	...				...		...				
?	?	1	5	U46	?	2	?	...	?	...	?	?	4	
?	?	1	7	...				...		...				
?	?	0	6	U98	5			...		...	4			
?	?	1	6	U99			1	...		...		2		


# Preference Learning Tasks

		representation		type of preference information		
		input	output	training	prediction	ground truth
generalized classification	collaborative filtering	identifier	<b>identifier</b>	absolute ordinal	absolute ordinal	absolute ordinal
	multilabel classification	feature	<b>identifier</b>	absolute binary	absolute binary	absolute binary
	multilabel ranking	feature	<b>identifier</b>	absolute binary	ranking	absolute binary
	graded multilabel classification	feature	<b>identifier</b>	absolute ordinal	absolute ordinal	absolute ordinal
	label ranking	feature	<b>identifier</b>	relative binary	ranking	ranking
	object ranking	<b>feature</b>	--	relative binary	ranking	ranking or subset
	instance ranking	<b>feature</b>	identifier	absolute ordinal	ranking	absolute ordinal

Two main directions: (1) ranking and variants (2) generalizations of classification.

# Loss Functions

## Things to be compared:



# References

---

- W. Cheng, K. Dembczynski and E. Hüllermeier. *Graded Multilabel Classification: The Ordinal Case*. ICML-2010, Haifa, Israel, 2010.
- W. Cheng and E. Hüllermeier. *Predicting partial orders: Ranking with abstention*. ECML/PKDD-2010, Barcelona, 2010.
- Y. Chevaleyre, F. Koriche, J. Lang, J. Mengin, B. Zanuttini. *Learning ordinal preferences on multiattribute domains: The case of CP-nets*. In: J. Fürnkranz and E. Hüllermeier (eds.) *Preference Learning*, Springer-Verlag, 2010.
- W.W. Cohen, R.E. Schapire and Y. Singer. *Learning to order things*. Journal of Artificial Intelligence Research, 10:243–270, 1999.
- J. Fürnkranz, E. Hüllermeier, E. Mencia, and K. Brinker. *Multilabel Classification via Calibrated Label Ranking*. Machine Learning 73(2):133-153, 2008.
- J. Fürnkranz, E. Hüllermeier and S. Vanderlooy. *Binary decomposition methods for multipartite ranking*. Proc. ECML-2009, Bled, Slovenia, 2009.
- D. Goldberg, D. Nichols, B.M. Oki and D. Terry. *Using collaborative filtering to weave and information tapestry*. Communications of the ACM, 35(12):61–70, 1992.
- E. Hüllermeier, J. Fürnkranz, W. Cheng and K. Brinker. *Label ranking by learning pairwise preferences*. Artificial Intelligence, 172:1897–1916, 2008.
- G. Tsoumakas and I. Katakis. *Multi-label classification: An overview*. Int. J. Data Warehouse and Mining, 3:1–13, 2007.
- A.K. Menon and C. Elkan. *Predicting labels for dyadic data*. Data Mining and Knowledge Discovery, 21(2), 2010



# AGENDA

---

1. Preference Learning Tasks
2. **Loss Functions**
  - a. **Evaluation of Rankings**
  - b. Weighted Measures
  - c. Evaluation of Bipartite Rankings
  - d. Evaluation of Partial Rankings
3. Preference Learning Techniques
4. Conclusions

# Rank Evaluation Measures

---

- In the following, we do not discriminate between different ranking scenarios
  - we use the term **items** for both, objects and labels
- All measures are applicable to both scenarios
  - sometimes have different names according to context
- Label Ranking
  - measure is applied to the ranking of the labels of each examples
  - averaged over all examples
- Object Ranking
  - measure is applied to the ranking of a set of objects
  - we may need to average over different sets of objects which have disjoint preference graphs
    - e.g. different sets of query / answer set pairs in information retrieval


# Ranking Errors

- Given:
  - a set of items  $X = \{x_1, \dots, x_c\}$  to rank

- Example:


$X = \{A, B, C, D, E\}$

items can be  
objects or labels




# Ranking Errors

- Given:
  - a set of items  $X = \{x_1, \dots, x_c\}$  to rank
    - *Example:*  
 $X = \{A, B, C, D, E\}$
  - a target ranking  $r$ 
    - *Example:*  
 $E > B > C > A > D$



# Ranking Errors

- Given:
  - a set of items  $X = \{x_1, \dots, x_c\}$  to rank
    - *Example:*  
 $X = \{A, B, C, D, E\}$
  - a target ranking  $r$ 
    - *Example:*  
 $E > B > C > A > D$
  - a predicted ranking  $\hat{r}$ 
    - *Example:*  
 $A > B > E > C > D$
- Compute:
  - a value  $d(r, \hat{r})$  that measures the *distance* between the two rankings



# Notation

- $r$  and  $\hat{r}$  are functions from  $X \rightarrow \mathbb{N}$ 
  - returning the rank of an item  $x$


$$\hat{r}(A) = 1$$

- the inverse functions  $r^{-1}: \mathbb{N} \rightarrow X$ 
  - return the item at a certain position

$$\hat{r}^{-1}(1) = A \quad r^{-1}(4) = A$$

- as a short-hand for  $r \circ \hat{r}^{-1}$ , we also define function  $R: \mathbb{N} \rightarrow \mathbb{N}$ 
  - $R(i)$  returns the true rank of the  $i$ -th item in the predicted ranking


$$R(1) = r(\hat{r}^{-1}(1)) = 4$$



# Spearman's Footrule

- Key idea:
  - Measure the sum of **absolute** differences between ranks

$$\begin{aligned} D_{SF}(r, \hat{r}) &= \sum_{i=1}^c |r(x_i) - \hat{r}(x_i)| = \sum_{i=1}^c |i - R(i)| \\ &= \sum_{i=1}^c d_{x_i}(r, \hat{r}) \end{aligned}$$



$$\sum_{x_i} d_{x_i} = 3 + 0 + 1 + 0 + 2 = 6$$

# Spearman Distance

- Key idea:
  - Measure the sum of ~~absolute~~ <sup>squared</sup> differences between ranks

$$D_S(r, \hat{r}) = \sum_{i=1}^c (r(x_i) - \hat{r}(x_i))^2 = \sum_{i=1}^c (i - R(i))^2$$
$$= \sum_{i=1}^c d_{x_i}(r, \hat{r})^2$$


- Value range:

$$\min D_S(r, \hat{r}) = 0$$

$$\max D_S(r, \hat{r}) = \sum_{i=1}^c ((c-i) - i)^2 = \frac{c \cdot (c^2 - 1)}{3}$$

→ Spearman Rank Correlation Coefficient

$$1 - \frac{6 \cdot D_S(r, \hat{r})}{c \cdot (c^2 - 1)} \in [-1, +1]$$



$$\sum_{x_i} d_{x_i}^2 = 3^2 + 0 + 1^2 + 0 + 2^2 = 14$$



# Kendall's Distance

- Key idea:
  - number of item pairs that are inverted in the predicted ranking

$$D_{\tau}(r, \hat{r}) = | \{ (i, j) \mid r(x_i) < r(x_j) \wedge \hat{r}(x_i) > \hat{r}(x_j) \} |$$


- Value range:

$$\min D_{\tau}(r, \hat{r}) = 0$$

$$\max D_{\tau}(r, \hat{r}) = \frac{c \cdot (c-1)}{2}$$

→ **Kendall's tau**

$$1 - \frac{4 \cdot D_{\tau}(r, \hat{r})}{c \cdot (c-1)} \in [-1, +1]$$



# AGENDA

---

1. Preference Learning Tasks
2. **Loss Functions**
  - a. Evaluation of Rankings
  - b. **Weighted Measures**
  - c. Evaluation of Bipartite Rankings
  - d. Evaluation of Partial Rankings
3. Preference Learning Techniques
4. Conclusions

# Weighted Ranking Errors

- The previous ranking functions give **equal weight to all ranking positions**
  - i.e., differences in the first ranking positions have the same effect as differences in the last ranking positions

$$D\left(\begin{array}{c} A \\ B \\ C \\ E \\ D \end{array}, \begin{array}{c} A \\ B \\ C \\ D \\ E \end{array}\right) = D\left(\begin{array}{c} A \\ B \\ C \\ D \\ E \end{array}, \begin{array}{c} B \\ A \\ C \\ D \\ E \end{array}\right)$$

- In many applications this is **not desirable**
  - ranking of search results
  - ranking of product recommendations
  - ranking of labels for classification
  - ...

⇒ Higher ranking positions should be given more weight

# Position Error


- Key idea:
  - in many applications we are interested in providing a ranking where the target item appears as high as possible in the predicted ranking
    - e.g. ranking a set of actions for the next step in a plan
  - Error is the number of wrong items that are predicted before the target item

$$D_{PE}(r, \hat{r}) = \hat{r}(\arg \min_{x \in X} r(x)) - 1$$

- Note:
  - equivalent to Spearman's footrule with all non-target weights set to 0

$$D_{PE}(r, \hat{r}) = \sum_{i=1}^c w_i \cdot d_{x_i}(r, \hat{r})$$

with  $w_i = \mathbb{1}_{x_i = \arg \min_{x \in X} r(x)}$




$$D_{PE}(r, \hat{r}) = 2$$

# Discounted Error

- Higher ranks in the target position get a higher weight than lower ranks

$$D_{DR}(r, \hat{r}) = \sum_{i=1}^c w_i \cdot d_{x_i}(r, \hat{r})$$

$$\text{with } w_i = \frac{1}{\log(r(x_i) + 1)}$$



$$D_{DR}(r, \hat{r}) = \frac{3}{\log 2} + 0 + \frac{1}{\log 4} + 0 + \frac{2}{\log 6}$$

# (Normalized) Discounted Cumulative Gain

- a “positive” version of discounted error:

## Discounted Cumulative Gain (DCG)


$$DCG(r, \hat{r}) = \sum_{i=1}^c \frac{c - R(i)}{\log(i+1)}$$

- Maximum possible value:
  - the predicted ranking is correct, i.e.  $\forall i: i = R(i)$
  - Ideal Discounted Cumulative Gain (IDCG)

$$IDCG = \sum_{i=1}^c \frac{c - i}{\log(i+1)}$$

- Normalized DCG (NDCG)

$$NDCG(r, \hat{r}) = \frac{DCG(r, \hat{r})}{IDCG}$$



$$NDCG(r, \hat{r}) = \frac{\frac{1}{\log 2} + \frac{3}{\log 3} + \frac{4}{\log 4} + \frac{2}{\log 5} + \frac{0}{\log 6}}{\frac{4}{\log 2} + \frac{3}{\log 3} + \frac{2}{\log 4} + \frac{1}{\log 5} + \frac{0}{\log 6}}$$

# AGENDA


---

1. Preference Learning Tasks
2. **Loss Functions**
  - a. Evaluation of Rankings
  - b. Weighted Measures
  - c. **Evaluation of Bipartite Rankings**
  - d. Evaluation of Partial Rankings
3. Preference Learning Techniques
4. Conclusions

# Bipartite Rankings

## Bipartite Rankings

- The target ranking is not totally ordered but a *bipartite graph*
- The two partitions may be viewed as preference levels  $L = \{0, 1\}$ 
  - all  $c_1$  items of level **1** are preferred over all  $c_0$  items of level **0**
- We now have fewer preferences
  - for a total order:  $\frac{c}{2} \cdot (c-1)$
  - for a bipartite graph:  $c_1 \cdot (c - c_1)$





# Evaluating Partial Target Rankings

- Many Measures can be directly adapted from total target rankings to partial target rankings

- Recall: **Kendall's distance**


- number of item pairs that are inverted in the target ranking

$$D_{\tau}(r, \hat{r}) = |\{(i, j) \mid r(x_i) < r(x_j) \wedge \hat{r}(x_i) > \hat{r}(x_j)\}|$$

- can be directly used
- in case of normalization, we have to consider that fewer items satisfy  $r(x_i) < r(x_j)$

- Area under the ROC curve (AUC)**

- the AUC is the fraction of pairs of  $(p, n)$  for which the predicted score  $s(p) > s(n)$ 
  - Mann Whitney statistic is the absolute number
- This is 1 - normalized Kendall's distance for a bipartite preference graph with  $L = \{p, n\}$



$$D_{\tau}(r, \hat{r}) = 2$$

$$AUC(r, \hat{r}) = \frac{4}{6}$$


# Evaluating Multipartite Rankings

- **Multipartite rankings:**
  - like Bipartite rankings
  - but the target ranking  $r$  consists of *multiple* relevance levels  $L = \{1 \dots l\}$ , where  $l < c$
  - total ranking is a special case where each level has exactly one item

- # of preferences  $= \sum_{(i,j)} c_i \cdot c_j \leq \frac{c^2}{2} \cdot (1 - \frac{1}{l})$ 
  - $c_i$  is the number of items in level  $i$

- **C-Index** [Gnen & Heller, 2005]
  - straight-forward generalization of AUC
  - fraction of pairs  $(x_i, x_j)$  for which

$$l(i) > l(j) \wedge \hat{r}(x_i) < \hat{r}(x_j)$$



$$D_{\tau}(r, \hat{r}) = 3$$

$$\text{C-Index}(r, \hat{r}) = \frac{5}{8}$$

# Evaluating Multipartite Rankings

## C-Index

- the C-index can be rewritten as a weighted sum of pairwise AUCs:

$$\text{C-Index}(r, \hat{r}) = \frac{1}{\sum_{i,j>i} c_i \cdot c_j} \sum_{i,j>i} c_i \cdot c_j \cdot \text{AUC}(r_{i,j}, \hat{r}_{i,j})$$

where  $r_{i,j}$  and  $\hat{r}_{i,j}$  are the rankings  $r$  and  $\hat{r}$  restricted to levels  $i$  and  $j$ .

## Jonckheere-Terpstra statistic

- is an *unweighted* sum of pairwise AUCs:

$$\text{m-AUC} = \frac{2}{l \cdot (l-1)} \sum_{i,j>i} \text{AUC}(r_{i,j}, \hat{r}_{i,j})$$

### Note:

C-Index and m-AUC can be optimized by optimization of pairwise AUCs

- equivalent to well-known multi-class extension of AUC [Hand & Till, MLJ 2001]


# Normalized Discounted Cumulative Gain

[Jarvelin & Kekalainen, 2002]

- The original formulation of (normalized) discounted cumulative gain refers to this setting

$$DCG(r, \hat{r}) = \sum_{i=1}^c \frac{l(i)}{\log(i+1)}$$

- the sum of the true (relevance) levels of the items
- each item weighted by its rank in the predicted ranking
- Examples:
  - retrieval of relevant or irrelevant pages
    - 2 relevance levels
  - movie recommendation
    - 5 relevance levels



# AGENDA

---

1. Preference Learning Tasks
2. **Loss Functions**
  - a. Evaluation of Rankings
  - b. Weighted Measures
  - c. Evaluation of Bipartite Rankings
  - d. **Evaluation of Partial Rankings**
3. Preference Learning Techniques
4. Conclusions

# Evaluating Partial Structures in the Predicted Ranking

---

- For fixed types of partial structures, we have conventional measures
  - bipartite graphs → binary classification
    - accuracy, recall, precision, F1, etc.
    - can also be used when the items are labels!
      - e.g., accuracy on the set of labels for multilabel classification
  - multipartite graphs → ordinal classification
    - multiclass classification measures (accuracy, error, etc.)
    - regression measures (sum of squared errors, etc.)
- For general partial structures
  - some measures can be directly used on the reduced set of target preferences
    - Kendall's distance, Gamma coefficient
  - we can also use **set measures** on the set of binary preferences
    - both, the source and the target ranking consist of a set of binary preferences
    - e.g. **Jaccard Coefficient**
      - size of intersection over size of union of the binary preferences in both sets

# Gamma Coefficient

- Key idea: normalized difference between

- number of **correctly** ranked pairs (Kendall's distance)

$$d = D_{\tau}(r, \hat{r})$$

- number of **incorrectly** ranked pairs

$$\bar{d} = |\{(i, j) \mid r(x_i) < r(x_j) \wedge \hat{r}(x_i) > \hat{r}(x_j)\}|$$


- **Gamma Coefficient**

[Goodman & Kruskal, 1979]

$$\gamma(r, \hat{r}) = \frac{d - \bar{d}}{d + \bar{d}} \in [-1, +1]$$

- Identical to Kendall's tau if both rankings are total

- i.e., if  $d + \bar{d} = \frac{c \cdot (c-1)}{2}$



$$\gamma(r, \hat{r}) = \frac{2-1}{2+1} = \frac{1}{3}$$

# References

---

- Cheng W., Rademaker M., De Baets B., Hüllermeier E.: *Predicting Partial Orders: Ranking with Abstention*. Proceedings ECML/PKDD-10(1): 215-230 (2010)
- Fürnkranz J., Hüllermeier E., Vanderlooy S.: *Binary Decomposition Methods for Multipartite Ranking*. Proceedings ECML/PKDD-09(1): 359-374 (2009)
- Gnen M., Heller G.: *Concordance probability and discriminatory power in proportional hazards regression*. Biometrika **92**(4):965–970 (2005)
- Goodman, L., Kruskal, W.: *Measures of Association for Cross Classifications*. Springer-Verlag, New York (1979)
- Hand D.J., Till R.J.: *A Simple Generalisation of the Area Under the ROC Curve for Multiple Class Classification Problems*. Machine Learning **45**(2):171-186 (2001)
- Jarvelin K., Kekalainen J.: *Cumulated gain-based evaluation of IR techniques*. ACM Transactions on Information Systems **20**(4): 422–446 (2002)
- Jonckheere, A. R.: *A distribution-free k-sample test against ordered alternatives*. Biometrika: 133–145 (1954)
- Kendall, M. *A New Measure of Rank Correlation*. Biometrika 30 (1-2): 81–89 (1938)
- Mann H. B., Whitney D. R. *On a test of whether one of two random variables is stochastically larger than the other*. Annals of Mathematical Statistics, **18**:50–60 (1947)
- Spearman C. *The proof and measurement of association between two things*. American Journal of Psychology, **15**:72–101 (1904)



# AGENDA

---

1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
- 3. Preference Learning Techniques**
  - a. Learning Utility Functions
  - b. Learning Preference Relations
  - c. Structured Output Prediction
  - d. Model-Based Preference Learning
  - e. Local Preference Aggregation
4. Conclusions

# Two Ways of Representing Preferences

- **Utility-based approach:** Evaluating single alternatives

$$U : \mathcal{A} \longrightarrow \mathbb{R}$$

- **Relational approach:** Comparing pairs of alternatives

$$a \succeq b \iff a \text{ is not worse than } b \quad \text{weak preference}$$

$$a \succ b \iff (a \succeq b) \wedge (b \not\succeq a) \quad \text{strict preference}$$

$$a \sim b \iff (a \succeq b) \wedge (b \succeq a) \quad \text{indifference}$$

$$a \perp b \iff (a \not\succeq b) \wedge (b \not\succeq a) \quad \text{incomparability}$$

# Utility Functions

- A **utility function** assigns a utility degree (typically a real number or an ordinal degree) to each alternative.
- Learning such a function essentially comes down to solving an (ordinal) **regression problem**.
- Often **additional conditions**, e.g., due to bounded utility ranges or monotonicity properties ( $\rightarrow$  *learning monotone models*)
- A **utility function induces a ranking** (total order), but not the other way around!
- But it can not represent more general relations, e.g., a **partial order**!
- The **feedback** can be **direct** (exemplary utility degrees given) or **indirect** (inequality induced by order relation):

$$(x, u) \Rightarrow U(x) \approx u, \quad x \succ y \Leftrightarrow U(x) > U(y)$$

absolute feedback

relative feedback

# Predicting Utilities on Ordinal Scales

(Graded) multilabel classification

X1	X2	X3	X4	A	B	C	D
0.34	0	10	174	--	+	++	0
1.45	0	32	277	0	++	--	+
1.22	1	46	421	--	--	0	+
0.74	1	25	165	0	+	+	++
0.95	1	72	273	+	0	++	--
1.04	0	33	158	+	+	++	--

Collaborative filtering

	P1	P2	P3	...	P38	...	P88	P89	P90
U1	1		4	...		...		3	
U2		2	2	...		...	1		
...				...		...			
U46	?	2	?	...	?	...	?	?	4
...				...		...			
U98	5			...		...	4		
U99			1	...		...		2	

Exploiting dependencies  
(correlations) between items  
(labels, products, ...)


→ see work in MLC and RecSys communities

# Learning Utility Functions from Indirect Feedback

- A (latent) utility function can also be used to solve ranking problems, such as instance, object or label ranking  
→ **ranking by (estimated) utility degrees (scores)**

## Object ranking

$(0.74, 1, 25, 165) \succ (0.45, 0, 35, 155)$   
 $(0.47, 1, 46, 183) \succ (0.57, 1, 61, 177)$   
 $(0.25, 0, 26, 199) \succ (0.73, 0, 46, 185)$   
 $(0.95, 0, 73, 133) \succ (0.25, 1, 35, 153)$   
 $(0.68, 1, 55, 147) \succ (0.67, 0, 63, 182)$




Find a utility function that agrees as much as possible with the preference information in the sense that, for most examples,

$$x_i \succ y_i \Leftrightarrow U(x_i) > U(y_i)$$

## Instance ranking


x1	x2	x3	x4	class
0.34	0	10	174	--
1.45	0	32	277	0
1.22	1	46	421	--
0.74	1	25	165	++
0.95	1	72	273	+




Absolute preferences given, so in principle an ordinal regression problem. However, the goal is to maximize **ranking** instead of **classification** performance.

# Ranking versus Classification

A ranker can be turned into a classifier via thresholding:



A good classifier is not necessarily a good ranker:



→ *learning **AUC-optimizing** scoring classifiers !*

## RankSVM and Related Methods (Bipartite Case)

- The idea is to minimize a convex upper bound on the empirical ranking error over a class of (kernelized) ranking functions:

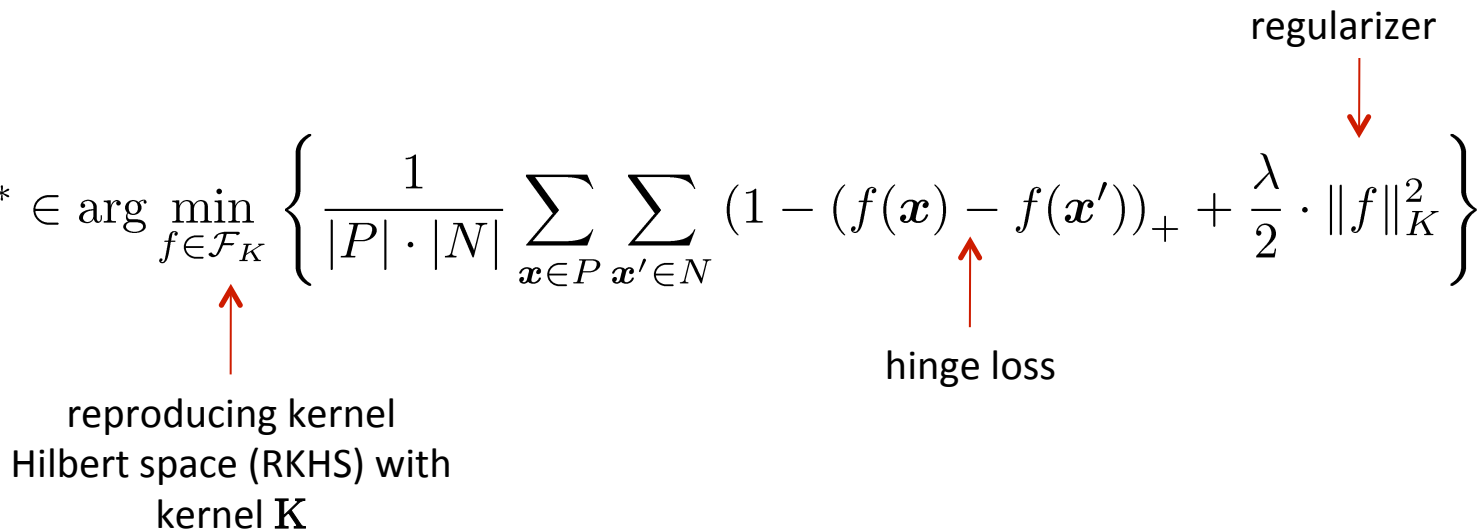
$$f^* \in \arg \min_{f \in \mathcal{F}} \left\{ \underbrace{\frac{1}{|P| \cdot |N|} \sum_{\mathbf{x} \in P} \sum_{\mathbf{x}' \in N} L(f, \mathbf{x}, \mathbf{x}')}_{\text{check for all positive/negative pairs}} + \lambda \cdot \underbrace{R(f)}_{\text{regularizer}} \right\}$$

convex upper bound on  $\mathbb{I}(f(\mathbf{x}) < f(\mathbf{x}'))$

## RankSVM and Related Methods (Bipartite Case)

- The bipartite RankSVM algorithm [Herbrich et al. 2000, Joachimes 2002]:

$$f^* \in \arg \min_{f \in \mathcal{F}_K} \left\{ \frac{1}{|P| \cdot |N|} \sum_{\mathbf{x} \in P} \sum_{\mathbf{x}' \in N} (1 - (f(\mathbf{x}) - f(\mathbf{x}'))_+) + \frac{\lambda}{2} \cdot \|f\|_K^2 \right\}$$



reproducing kernel  
Hilbert space (RKHS) with  
kernel  $K$

hinge loss

regularizer

→ learning comes down to solving a QP problem



## RankSVM and Related Methods (Bipartite Case)

- The bipartite RankBoost algorithm [Freund et al. 2003]:

$$f^* \in \arg \min_{f \in \mathcal{L}(\mathcal{F}_{base})} \left\{ \frac{1}{|P| \cdot |N|} \sum_{\mathbf{x} \in P} \sum_{\mathbf{x}' \in N} \exp(-(f(\mathbf{x}) - f(\mathbf{x}')))) \right\}$$

↑  
class of linear  
combinations of base  
functions

→ learning by means of boosting techniques

# Learning Utility Functions for Label Ranking

Label ranking is the problem of learning a function  $\mathcal{X} \rightarrow \Omega$ , with  $\Omega$  the set of rankings (permutations) of a label set  $\mathcal{Y} = \{y_1, y_2, \dots, y_k\}$ , from exemplary pairwise preferences  $y_i \succ_x y_j$ .

Can be tackled by learning utility functions  $U_1(\cdot), \dots, U_k(\cdot)$  that are in appropriate agreement with the preferences in the training data. Given a new query  $x$ , the labels are ranked according to utility degrees, i.e., a permutation  $\pi$  is predicted such that

$$U_{\pi^{-1}(1)}(x) > U_{\pi^{-1}(2)}(x) > \dots > U_{\pi^{-1}(k)}(x)$$

## Label Ranking: Reduction to Binary Classification [Har-Peled et al. 2002]

Proceeding from linear utility functions

$$U_i(\mathbf{x}) = \mathbf{w}_i \times \mathbf{x} = (w_{i,1}, w_{i,2}, \dots, w_{i,m})(x_1, x_2, \dots, x_m)^\top,$$

a binary preference  $y_i \succ_x y_j$  is equivalent to

$$U_i(\mathbf{x}) > U_j(\mathbf{x}) \Leftrightarrow \mathbf{w}_i \times \mathbf{x} > \mathbf{w}_j \times \mathbf{x} \Leftrightarrow (\mathbf{w}_i - \mathbf{w}_j) \times \mathbf{x} > 0$$

and can be modeled as a linear constraint

$$(\mathbf{w}_1, \mathbf{w}_2 \dots \mathbf{w}_k) \times (0 \dots 0 \mathbf{x} 0 \dots 0 - \mathbf{x} 0 \dots 0)^\top > 0$$

$(m \times k)$ -dimensional weight vector    positive example in the new instance space

Each **pairwise comparison** is turned into a **binary classification** example in a high-dimensional space!

# AGENDA

---

1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
3. Preference Learning Techniques
  - a. Learning Utility Functions
  - b. Learning Preference Relations**
  - c. Structured Output Prediction
  - d. Model-Based Preference Learning
  - e. Local Preference Aggregation
4. Conclusions

# Learning Binary Preference Relations

- Learning **binary preferences** (in the form of predicates  $P(x,y)$ ) is often simpler, especially if the training information is given in this form, too.
- However, it implies an additional step, namely **extracting a ranking** from a (predicted) preference relation.
- This step is not always trivial, since a predicted preference relation may exhibit inconsistencies and may not suggest a unique ranking in an unequivocal way.

instance $x$	$\longrightarrow$	$f_{i,j}$	$y_1$	$y_2$	$y_3$	$y_4$	$y_5$	$\xrightarrow{\text{inference}}$	$y_4 \succ y_5 \succ y_1 \succ y_3 \succ y_2$
		$y_1$		1	1	0	0		
		$y_2$	0		0	1	0		
		$y_3$	0	1		0	0		
		$y_4$	1	0	1		1		
		$y_5$	1	1	1	0			

## Object Ranking: Learning to Order Things [Cohen et al. 99]

- In a first step, a **binary preference function** **PREF** is constructed; **PREF**  $(\mathbf{x}, \mathbf{y}) \in [0, 1]$  is a measure of the certainty that  $\mathbf{x}$  should be ranked before  $\mathbf{y}$ , and  $\text{PREF}(\mathbf{x}, \mathbf{y}) = 1 - \text{PREF}(\mathbf{y}, \mathbf{x})$ .
- This function is expressed as a linear combination of base preference functions:

$$\text{PREF}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^N w_i \cdot R_i(\mathbf{x}, \mathbf{y})$$



- The weights can be learned, e.g., by means of the weighted majority algorithm [Littlestone & Warmuth 94].
- In a second step, a total order is derived, which is as much as possible in agreement with the binary preference relation.

# Object Ranking: Learning to Order Things [Cohen et al. 99]

- The weighted feedback arc set problem: Find a permutation  $\pi$  such that

$$\sum_{(x,y): \pi(x) > \pi(y)} \text{PREF}(x, y)$$

becomes minimal.



$$\text{cost} = 0.1 + 0.6 + 0.8 + 0.5 + 0.3 + 0.4 = 2.7$$

## Object Ranking: Learning to Order Things [Cohen et al. 99]

- Since this is an NP-hard problem, it is solved heuristically.

```

Input: an instance set  $X$ ; a preference function  $\text{PREF}$ 
Output: an approximately optimal ordering function  $\hat{\rho}$ 
let  $V = X$ 
for each  $v \in V$  do
  while  $V$  is non-empty do  $\pi(v) = \sum_{u \in V} \text{PREF}(v, u) - \sum_{u \in V} \text{PREF}(u, v)$ 
    let  $t = \arg \max_{u \in V} \pi(u)$ 
    let  $\hat{\rho}(t) = |V|$ 
     $V = V - \{t\}$ 
    for each  $v \in V$  do  $\pi(v) = \pi(v) + \text{PREF}(t, v) - \text{PREF}(v, t)$ 
  endwhile
  
```

- The algorithm successively chooses nodes having **maximal „net-flow“** within the remaining subgraph.
- It can be shown to provide a **2-approximation** to the optimal solution.



## Label Ranking: Learning by Pairwise Comparison (LPC) [Hüllermeier et al. 2008]

Label ranking is the problem of learning a function  $\mathcal{X} \rightarrow \Omega$ , with  $\Omega$  the set of rankings (permutations) of a label set  $\mathcal{Y} = \{y_1, y_2, \dots, y_k\}$ , from exemplary pairwise preferences  $y_i \succ_x y_j$ .

LPC trains a model

$$\mathcal{M}_{i,j} : \mathcal{X} \rightarrow [0, 1]$$

for all  $i < j$ . Given a query instance  $x$ , this model is supposed to predict whether  $y_i \succ y_j$  ( $\mathcal{M}_{i,j}(x) = 1$ ) or  $y_j \succ y_i$  ( $\mathcal{M}_{i,j}(x) = 0$ ).

More generally,  $\mathcal{M}_{i,j}(x)$  is the estimated probability that  $y_i \succ y_j$ .

Decomposition into  $k(k-1)/2$  **binary classification problems**.

# Label Ranking: Learning by Pairwise Comparison (LPC) [Hüllermeier et al. 2008]

Training data (for the label pair A and B):

X1	X2	X3	X1	X2	X3	X4	class	class
0.34	0	10	0.34	0	10	174	1	1
1.45	0	32	1.22	1	46	421	0	0
1.22	1	46	0.74	1	25	165	1	1
0.74	1	25	1.04	0	33	158	1	1
0.95	1	72	273	B $\succ$ D, A $\succ$ D,				
1.04	0	33	158	D $\succ$ A, A $\succ$ B, C $\succ$ B, A $\succ$ C				1

## Label Ranking: Learning by Pairwise Comparison (LPC) [Hüllermeier et al. 2008]

At prediction time, a query instance is submitted to all models, and the predictions are combined into a binary preference relation:

predictions  
 $\mathcal{M}_{i,j}(\mathbf{x})$  →

	A	B	C	D
A		0.3	0.8	0.4
B	0.7		0.7	0.9
C	0.2	0.3		0.3
D	0.6	0.1	0.7	

## Label Ranking: Learning by Pairwise Comparison (LPC) [Hüllermeier et al. 2008]

At prediction time, a query instance is submitted to all models, and the predictions are combined into a binary preference relation:

predictions  $\mathcal{M}_{i,j}(\mathbf{x}) \longrightarrow$

	A	B	C	D	
A		0.3	0.8	0.4	1.5
B	0.7		0.7	0.9	2.3
C	0.2	0.3		0.3	0.8
D	0.6	0.1	0.7		1.4

$B \succ A \succ D \succ C$

From this relation, a ranking is derived by means of a **ranking procedure**. In the simplest case, this is done by sorting the labels according to their sum of **weighted votes**.

---

# AGENDA

---

1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
3. Preference Learning Techniques
  - a. Learning Utility Functions
  - b. Learning Preference Relations
  - c. Structured Output Prediction**
  - d. Model-Based Preference Learning
  - e. Local Preference Aggregation
4. Conclusions

## Structured Output Prediction [Bakir et al. 2007]

- Rankings, multilabel classifications, etc. can be seen as specific types of **structured** (as opposed to scalar) **outputs**.
- Discriminative structured prediction algorithms infer a **joint scoring function on input-output pairs** and, for a given input, predict the output that maximises this scoring function.
- Joint feature map and scoring function

$$\phi : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}^d, \quad f(\mathbf{x}, \mathbf{y}; \mathbf{w}) = \langle \mathbf{w}, \phi(\mathbf{x}, \mathbf{y}) \rangle$$

- The learning problem consists of estimating the weight vector, e.g., using structural risk minimization.
- Prediction requires solving a **decoding problem**:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}} f(\mathbf{x}, \mathbf{y}; \mathbf{w}) = \arg \max_{\mathbf{y} \in \mathcal{Y}} \langle \mathbf{w}, \phi(\mathbf{x}, \mathbf{y}) \rangle$$

# Structured Output Prediction [Bakir et al. 2007]

- **Preferences** are expressed through **inequalities** on inner products:

$$\min_{\mathbf{w}, \xi} ||\mathbf{w}||^2 + \nu \sum_{i=1}^m \xi_i$$

loss function  
↓

$$\text{s.t. } \langle \mathbf{w}, \phi(\mathbf{x}_i, \mathbf{y}_i) \rangle - \langle \mathbf{w}, \phi(\mathbf{x}_i, \mathbf{y}) \rangle \geq \Delta(\mathbf{y}_i, \mathbf{y}) - \xi_i \text{ for all } \mathbf{y} \in \mathcal{Y}$$

$$\xi_i \geq 0 \quad (i = 1, \dots, m)$$

- The potentially huge **number of constraints** cannot be handled explicitly and calls for specific techniques (such as cutting plane optimization)

# AGENDA

---

1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
3. Preference Learning Techniques
  - a. Learning Utility Functions
  - b. Learning Preference Relations
  - c. Structured Output Prediction
  - d. Model-Based Preference Learning**
  - e. Local Preference Aggregation
4. Conclusions



# Model-Based Methods for Ranking

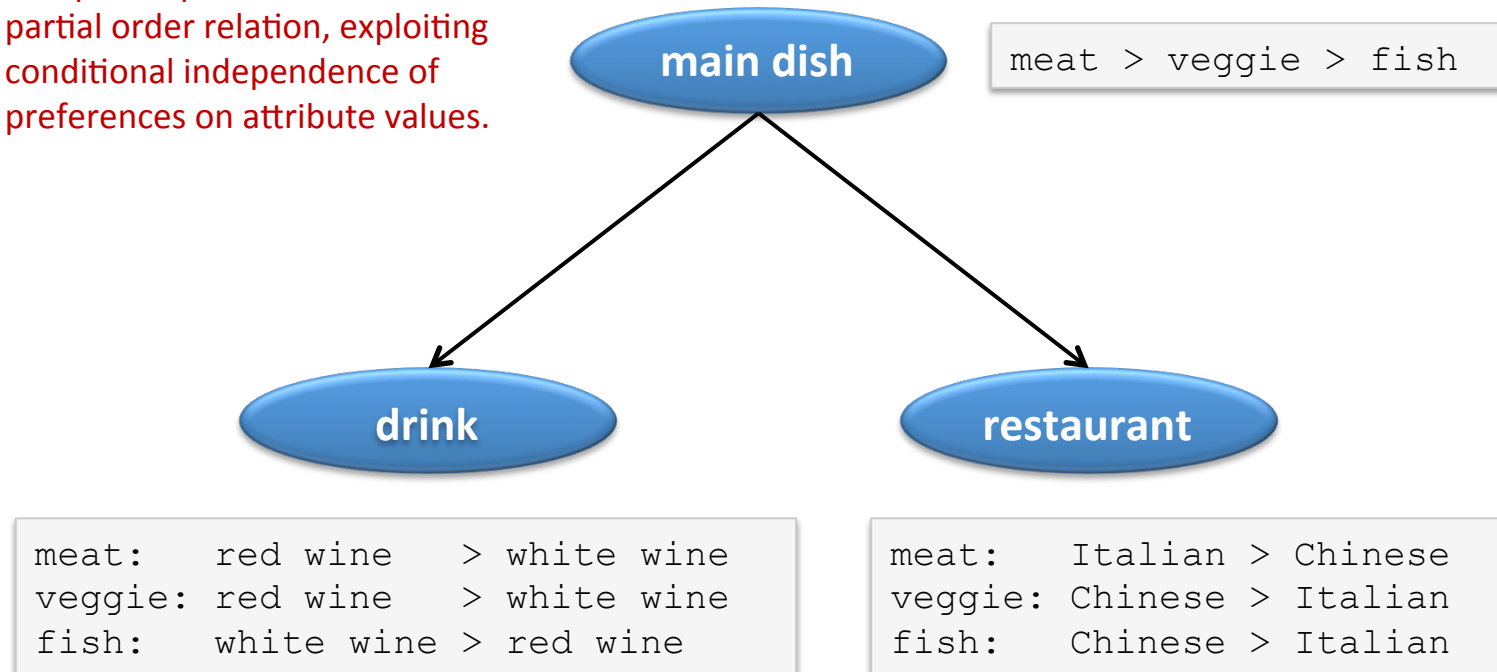
- **Model-based approaches** to ranking proceed from specific assumptions about the possible rankings (**representation bias**) or make use of **probabilistic models** for rankings (parametrized probability distributions on the set of rankings).
- In the following, we shall see examples of both type:
  - Restriction to lexicographic preferences
  - Conditional preference networks (CP-nets)
  - ~~Label ranking using the Plackett-Luce model~~ **see our talk tomorrow**

# Learning Lexicographic Preference Models [Yaman et al. 2008]

- Suppose that objects are represented as feature vectors of length  $m$ , and that each attribute has  $k$  values.
- For  $n=k^m$  objects, there are  $n!$  permutations (rankings).
- A **lexicographic order** is uniquely determined by
  - a total order of the attributes
  - a total order of each attribute domain
- **Example:** Four binary attributes ( $m=4$ ,  $k=2$ )
  - there are  $16! \approx 2 \cdot 10^{13}$  rankings
  - but only  $(2^4) \cdot 4! = 384$  of them can be expressed in terms of a lexicographic order
- [Yaman et al. 2008] present a learning algorithm that explicitly maintains the version space, i.e., the attribute-orders compatible with all pairwise preferences seen so far (assuming binary attributes with 1 preferred to 0). Predictions are derived based on the „votes“ of the consistent models.

# Learning Conditional Preference (CP) Networks [Chevaleyre et al. 2010]

Compact representation of a partial order relation, exploiting conditional independence of preferences on attribute values.

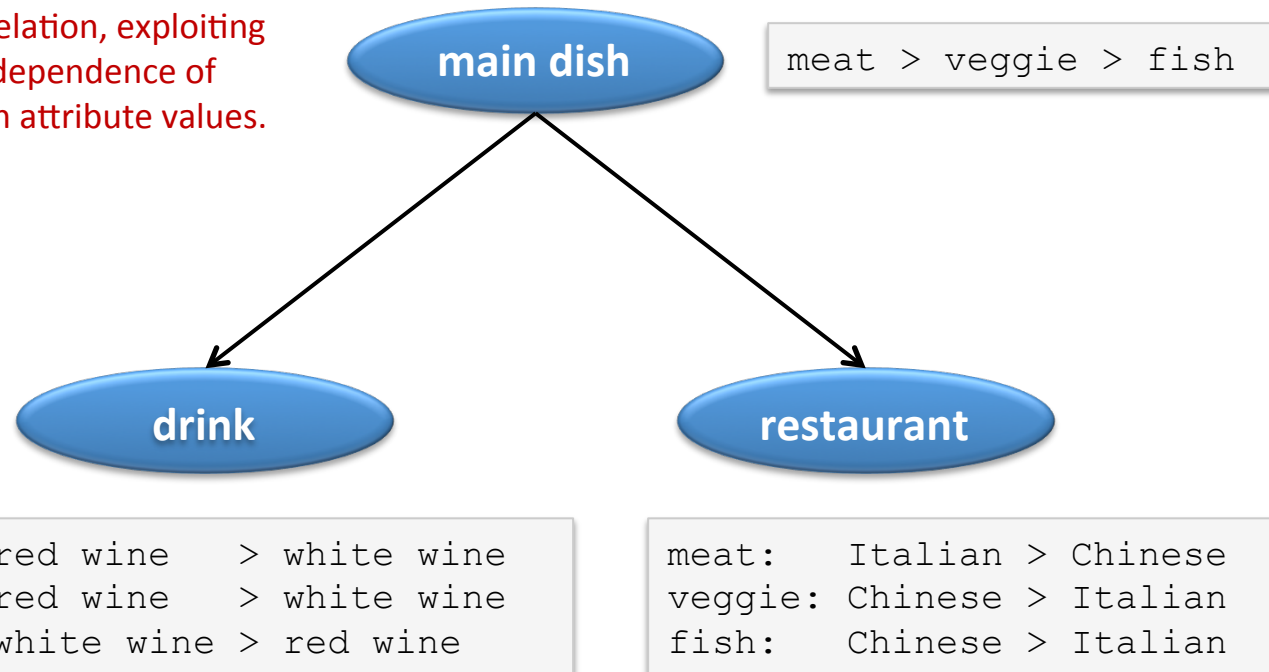


Induces partial order relation, e.g.,

```
(meat, red wine, Italian) > (meat, white wine, Chinese)
(fish, white wine, Chinese) > (fish, red wine, Chinese)
(meat, white wine, Italian) ? (meat, red wine, Chinese)
```

# Learning Conditional Preference (CP) Networks [Chevaleyre et al. 2010]

Compact representation of a partial order relation, exploiting conditional independence of preferences on attribute values.




**Training data (possibly noisy):**

```
(meat, red wine, Italian) > (veggie, red wine, Italian)
(fish, whited wine, Chinease) > (veggie, red wine, Chinease)
(veggie, whited wine, Chinease) > (veggie, red wine, Italian)
... ..
```

# AGENDA

1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
3. Preference Learning Techniques
  - a. Learning Utility Functions
  - b. Learning Preference Relations
  - c. Structured Output Prediction
  - d. Model-Based Preference Learning
  - e. Local Preference Aggregation **see our talk tomorrow**
4. Conclusions



## Summary of Main Algorithmic Principles

---

- **Reduction** of ranking to (binary) classification (e.g., constraint classification, LPC)
- **Direct optimization** of (regularized) smooth approximation of ranking losses (RankSVM, RankBoost, ...)
- **Structured output prediction**, learning joint scoring („matching“) function
- Learning parametrized **probabilistic ranking models** (e.g., Plackett-Luce)
- **Restricted model classes**, fitting (parametrized) deterministic models (e.g., lexicographic orders)
- **Lazy learning**, local preference aggregation (lazy learning)

# References

- G. Bakir, T. Hofmann, B. Schölkopf, A. Smola, B. Taskar and S. Vishwanathan. *Predicting structured data*. MIT Press, 2007.
- W. Cheng, K. Dembczynski and E. Hüllermeier. *Graded Multilabel Classification: The Ordinal Case*. ICML-2010, Haifa, Israel, 2010.
- W. Cheng, K. Dembczynski and E. Hüllermeier. *Label ranking using the Plackett-Luce model*. ICML-2010, Haifa, Israel, 2010.
- W. Cheng and E. Hüllermeier. *Predicting partial orders: Ranking with abstention*. ECML/PKDD-2010, Barcelona, 2010.
- W. Cheng, C. Hühn and E. Hüllermeier. *Decision tree and instance-based learning for label ranking*. ICML-2009.
- Y. Chevaleyre, F. Koriche, J. Lang, J. Mengin, B. Zanuttini. *Learning ordinal preferences on multiattribute domains: The case of CP-nets*. In: J. Fürnkranz and E. Hüllermeier (eds.) *Preference Learning*, Springer-Verlag, 2010.
- W.W. Cohen, R.E. Schapire and Y. Singer. *Learning to order things*. Journal of Artificial Intelligence Research, 10:243–270, 1999.
- Y. Freund, R. Iyer, R. E. Schapire and Y. Singer. *An efficient boosting algorithm for combining preferences*. Journal of Machine Learning Research, 4:933–969, 2003.
- J. Fürnkranz, E. Hüllermeier, E. Mencia, and K. Brinker. *Multilabel Classification via Calibrated Label Ranking*. Machine Learning 73(2):133-153, 2008.
- J. Fürnkranz, E. Hüllermeier and S. Vanderlooy. *Binary decomposition methods for multipartite ranking*. Proc. ECML-2009, Bled, Slovenia, 2009.
- D. Goldberg, D. Nichols, B.M. Oki and D. Terry. *Using collaborative filtering to weave and information tapestry*. Communications of the ACM, 35 (12):61–70, 1992.
- S. Har-Peled, D. Roth and D. Zimak. *Constraint classification: A new approach to multiclass classification*. Proc. ALT-2002.
- R. Herbrich, T. Graepel and K. Obermayer. *Large margin rank boundaries for ordinal regression*. Advances in Large Margin Classifiers, 2000.
- E. Hüllermeier, J. Fürnkranz, W. Cheng and K. Brinker. *Label ranking by learning pairwise preferences*. Artificial Intelligence, 172:1897–1916, 2008.
- T. Joachims. *Optimizing search engines using clickthrough data*. Proc. KDD 2002.
- N. Littlestone and M.K. Warmuth. *The weighted majority algorithm*. Information and Computation, 108(2): 212–261, 1994.
- G. Tsoumakas and I. Katakis. *Multilabel classification: An overview*. Int. J. Data Warehouse and Mining, 3:1–13, 2007.
- F. Yaman, T. Walsh, M. Littman and M. desJardins. *Democratic Approximation of Lexicographic Preference Models*. ICML-2008.

---

# AGENDA

---

1. Preference Learning Tasks
2. Performance Assessment and Loss Functions
3. Preference Learning Techniques
- 4. Conclusions**



# Conclusions

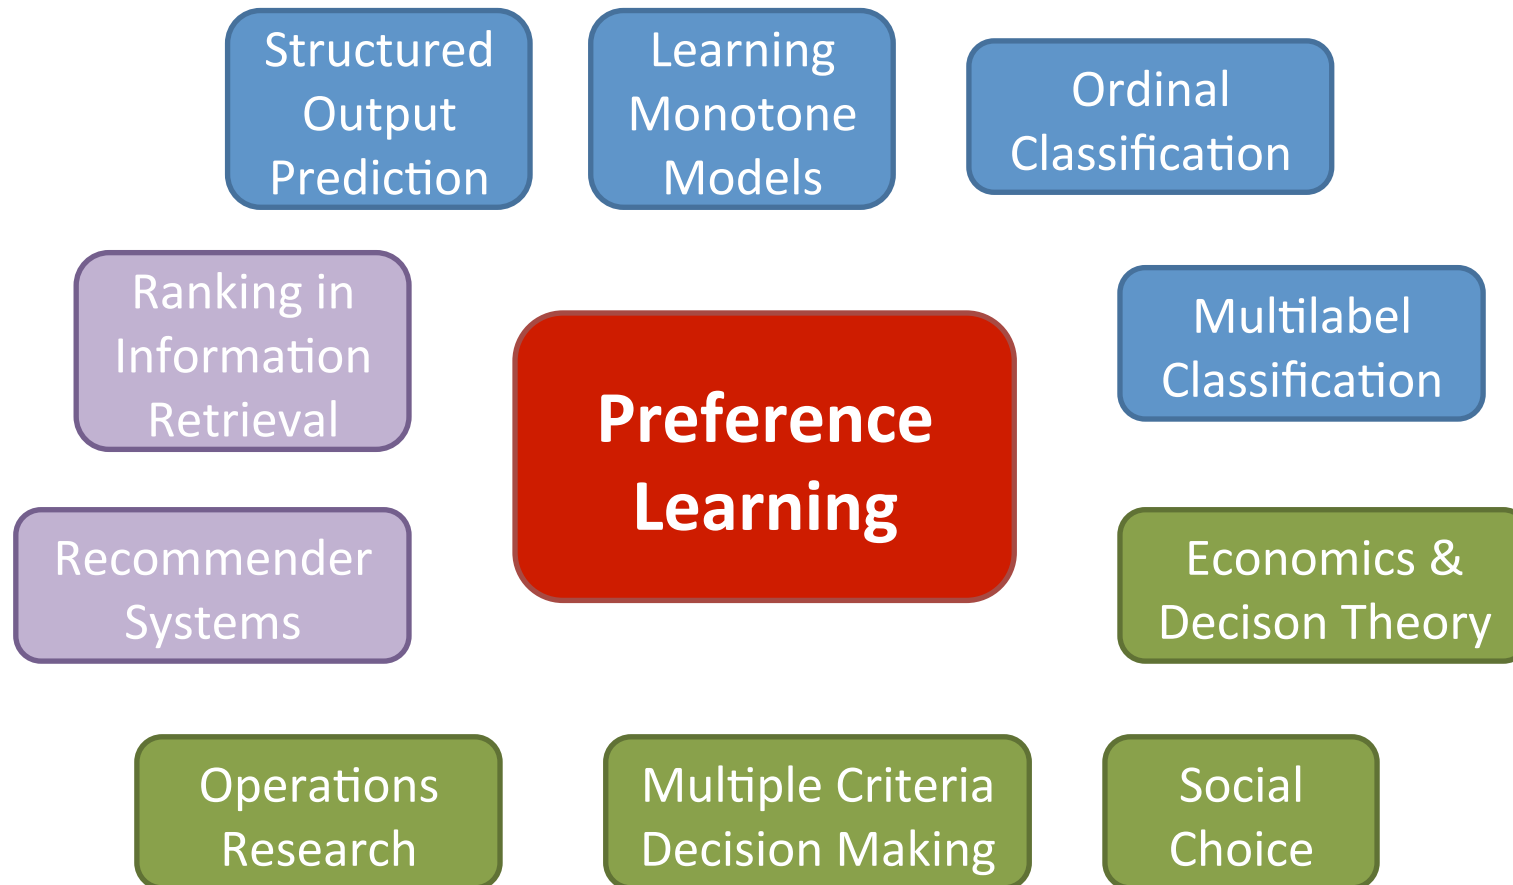
---

- Preference learning is an **emerging subfield** of machine learning, with many **applications** and **theoretical challenges**.
- Prediction of **preference models** instead of scalar outputs (like in classification and regression), hitherto with a focus on **rankings**.
- Many existing machine learning problems can be cast in the framework of preference learning (→ preference learning „in a broad sense“)
- **„Qualitative“ alternative** to conventional numerical approaches
  - pairwise comparison instead of numerical evaluation,
  - order relations instead of individual assessment.
- Still many **open problems** (unified framework, predictions more general than rankings, incorporating numerical information, etc.)
- **Interdisciplinary field**, connections to many other areas.

---

## Connections to Other Fields

---



---

# Edited Book on Preference Learning

---

Preference Learning: An Introduction

A Preference Optimization based Unifying Framework for Supervised Learning Problems

## **Part I – Label Ranking**

Label Ranking Algorithms: A Survey

Preference Learning and Ranking by Pairwise Comparison

Decision Tree Modeling for Ranking Data

Co-regularized Least-Squares for Label Ranking

## **Part II – Instance Ranking**

A Survey on ROC-Based Ordinal Regression

Ranking Cases with Classification Rules

## **Part III – Object Ranking**

A Survey and Empirical Comparison of Object Ranking Methods

Dimension Reduction for Object Ranking

Learning of Rule Ensembles for Multiple Attribute Ranking Problems

## **Part IV – Preferences in Multiattribute Domains**

Learning Lexicographic Preference Models

Learning Ordinal Preferences on Multiattribute Domains: the Case of CP-nets

Choice-Based Conjoint Analysis: Classification vs. Discrete Choice Models

Learning Aggregation Operators for Preference Modeling

## **Part V – Preferences in Information Retrieval**

Evaluating Search Engine Relevance with Click-Based Metrics


Learning SVM Ranking Function from User Feedback Using Document Metadata and Active Learning in the Biomedical Domain

## **Part VI – Preferences in Recommender Systems**

Learning Preference Models in Recommender Systems

Collaborative Preference Learning

Discerning Relevant Model Features in a Content-Based Collaborative Recommender System



J. Fürnkranz &  
E. Hüllermeier (eds.)  
Preference Learning  
Springer-Verlag 2010

# Edited Book on Preference Learning

## **Preference Learning: An Introduction**

A Preference Optimization based Unifying Framework for Supervised Learning Problems

### **Part I – Label Ranking**

#### **Label Ranking Algorithms: A Survey**

Preference Learning and Ranking by Pairwise Comparison

Decision Tree Modeling for Ranking Data

Co-regularized Least-Squares for Label Ranking

### **Part II – Instance Ranking**

#### **A Survey on ROC-Based Ordinal Regression**

Ranking Cases with Classification Rules

### **Part III – Object Ranking**

#### **A Survey and Empirical Comparison of Object Ranking Methods**

Dimension Reduction for Object Ranking

Learning of Rule Ensembles for Multiple Attribute Ranking Problems

### **Part IV – Preferences in Multiattribute Domains**

Learning Lexicographic Preference Models

Learning Ordinal Preferences on Multiattribute Domains: the Case of CP-nets

Choice-Based Conjoint Analysis: Classification vs. Discrete Choice Models

Learning Aggregation Operators for Preference Modeling

### **Part V – Preferences in Information Retrieval**

Evaluating Search Engine Relevance with Click-Based Metrics

Learning SVM Ranking Function from User Feedback Using Document Metadata and Active Learning in the Biomedical Domain


### **Part VI – Preferences in Recommender Systems**

#### **Learning Preference Models in Recommender Systems**

Collaborative Preference Learning

Discerning Relevant Model Features in a Content-Based Collaborative Recommender System

includes several introductions  
and survey articles



J. Fürnkranz &  
E. Hüllermeier (eds.)  
Preference Learning  
Springer-Verlag 2010

---

## Preference Learning Website

---

<http://www.preference-learning.org/>

- Working groups
- Software
- Data Sets
- Workshops
- Tutorials
- Books
- ...